

# Multi-Agent Path Finding

Luigi Palopoli, Enrico Saccon




**UNIVERSITÀ  
DI TRENTO**

# What?

- The standard Multi-Agent Path Finding (MAPF) problem [1] consists in:  
given a map and  $N$  agents, finding the best *feasible joint plan*  $\Pi$  such that each agent moves from its initial position to its final position minimizing an *objective function*
- It's a combinatorial problem
- In the standard definition, the map is usually a **grid**
- Solution minimizes objective function --> time, space, resource, etc.
- **Centralized approach**


# Why?

- Robotics is a main topic in the Industrial Revolution 4.0 and 5.0
- Used in an increasing number of scenarios
- Challenging problem [2]
- Algorithms can be applied to other scenarios




# Map Decomposition

- How do I deal with obstacles and map borders?
- Many algorithms to partition the map in cells:
  - Exact cell decomposition
  - Approximate cell decomposition
  - Maximum clearance
  - Morse decomposition
  - Brushfire decomposition
- Each cell is a node --> connectivity graph




# Map Decomposition

- How do I deal with obstacles and map borders?
- Many algorithms to partition the map in cells:
  - Exact cell decomposition
  - Approximate cell decomposition
  - Maximum clearance
  - Morse decomposition
  - Brushfire decomposition
- Each cell is a node --> connectivity graph



# Single-Agent Path Finding (SAPF)

- Given a graph  $G = (V, E)$ , find the *best feasible* plan  $\pi_i$  to go from an initial position to a final position
- The problem usually consists in computing the shortest path between two nodes on a graph
- Deterministic algorithms, e.g., Dijkstra's
- Heuristic algorithms, e.g.,  $A^*$




# Single-Agent Path Finding (SAPF)

## Dijkstra's

- Complete
- Optimal
- Evolution of BFS
- Cost function:

$$f(x) = g(x)$$




## A\*

- Complete?
- Optimal?
- Admissible heuristic  $h(x)$ :


$$f(x) = g(x) + h(x)$$

$$h(x) \leq d(x, y) + h(y)$$




# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- A joint plan is a set of single plans:  $\Pi = \{\pi_1, \dots, \pi_k\}$
- A path is feasible if no conflict arises [1]:
  - Vertex conflicts
  - Edge conflicts
  - Swap conflicts
- Objective functions:
  - Makespan (MKS)
  - Sum of individual costs (SIC)



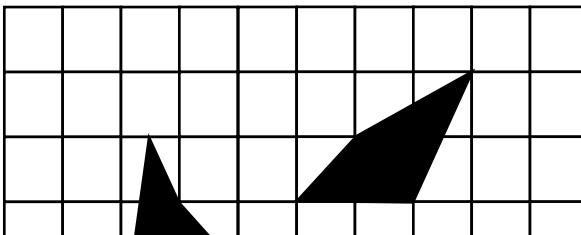




# Multi-Agent Path Finding Cost Functions



$\Pi_i$	$\text{SIC}(\Pi_i)$	$\text{MKS}(\Pi_i)$
$\Pi_1 = \begin{cases} \pi_1 = \{3, 10, 11, 12, 13, 7\} \\ \pi_2 = \{2, 4, 8, 9\} \\ \pi_3 = \{1, 2, 4, 8\} \end{cases}$	14	6
$\Pi_2 = \begin{cases} \pi_1 = \{3, 4, 5, 6, 7\} \\ \pi_2 = \{2, 2, 4, 8, 9\} \\ \pi_3 = \{1, 1, 2, 4, 8\} \end{cases}$	15	5



# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
  - Time is **discretized**
  - Each agent can either
    - move to an adjacent cell; or
    - stay on the same cell --> variation
- 





# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs





# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs





# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs





# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs





# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs



# Multi-Agent Path Finding Overview



- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs





# Multi-Agent Path Finding Overview

- Given a graph  $G = (V, E)$  and  $k$  agents, find the *best feasible joint* plan  $\Pi$  such that each agent moves from its initial position to its final position minimizing an objective function
- Time is **discretized**
- Each agent can either
  - move to an adjacent cell; or
  - stay on the same cell --> variation
- Edges have **unitary** costs


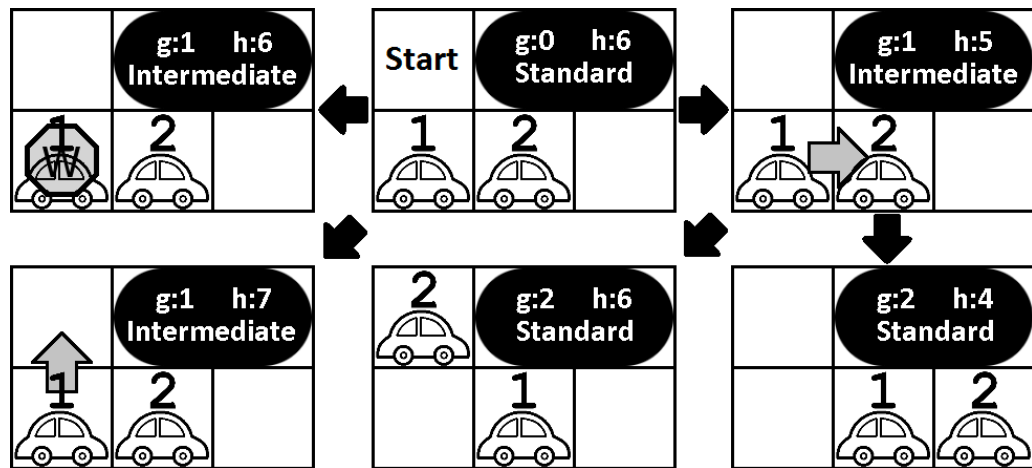


# Enhanced Versions of A\* [3]

- Instead of considering one position, considers a tuple
- At each timestep:
  - the current state space contains the position of all  $N$  agents
  - the next state space has to consider all possible movements of the agents
- This gets bad pretty fast
- **Operator decomposition (OD)**
- **Simple independence detection (SID)**

# Enhanced Versions of $A^*$ – OD

- Do not consider  $N$  agents per each time step
- Considers 1 agent at a time and it requires  $N$  operations to advance 1 timestep
- The order in which the agents are chosen is fixed
- It's not complete or optimal without the correct heuristic



# Enhanced Versions of A\* – SID



- Also in OD, the search space is still exponential in the number of agents
- Agents whose path *does not collide*, are in **independent** group and should be considered separately
- The algorithm follow these steps:
  - Start with the optimal paths as if the agents were alone
  - If there are conflicts, divide the agents in conflict groups
  - Solve the conflicts in the group
  - Repeat

# Priority Planning (PP) [4]

- Each agent has a fixed priority
- The higher the priority, the first the agent's path is computed
- Pros:
  - This allows for keeping the complexity small
- Cons:
  - Not complete [5]
  - Not optimal
- Many works, focus on using ML or DL approaches to learning the priorities


# Priority Planning (PP)

- There are some instances in which priority planning cannot solve the problem [5]




# Priority Planning (PP)

- There are some instances in which priority planning cannot solve the problem [5]



Not solvable



Solvable

- A MAPF instance must be *well-formed* to be solvable
  - Agents can wait for any amount of time on the initial or final node without blocking other agents

# Conflict Based Search (CBS)

- Proposed in 2015 by G. Sharon, R. Stern, A. Felner, and R. Holte [6]
- Creates a *constraint tree*

- Optimal algorithm divided in two phases:
  - When a conflict is found, it creates

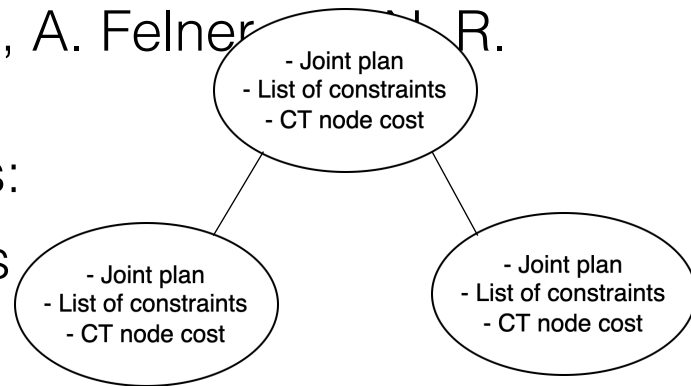
1. High-level search → manages conflicts

2. Low-level search → SAPF problem

- Agent  $a_i$  cannot be on node  $n$  at time  $t$
- Agent  $a_j$  cannot be on node  $n$  at time  $t$
- The search continues until a joint plan without conflicts is found
- Nodes to be explored are chosen based on their joint cost

2. Low-level search → SAPF problem

- The algorithm should be adapted to the problem





# CBS Implementation – High-Level

- Start from a root node with:
  - No constraints
  - Joint plan computed as SAPF
- Iterate until a feasible solution is found
  - If one or more vertex conflicts were found, then create two new nodes:
    - Child 1: agent  $a_i$  cannot be on node  $n$  at time  $t$
    - Child 2: agent  $a_j$  cannot be on node  $n$  at time  $t$
  - If one or more swap conflicts were found, then create two new nodes:
    - Child 1: agent  $a_i$  cannot move from node  $n_1$  to node  $n_2$  at time  $t$
    - Child 2: agent  $a_j$  cannot move from node  $n_2$  to node  $n_1$  at time  $t$
- Conflicts are checked by comparing the positions of the solutions
- A joint plan is updated only for the new constraint's agent

# CBS Implementation – Low-Level Spanning Tree


- Observation: difficult to find alternative paths when faced with constraints
- The algorithm computes all the possible paths between two points on a graph
- Then it starts from the shortest path and check if any conflicts arises with the constraints
- If it does, then it inserts waiting actions
- Finally, the shortest path is returned

# CBS Implementation – Low-Level TDSP

- This algorithm strongly takes from Dijkstra's
- The connectivity matrix is changed with Connection types
- Connection stores:
  - A vector of time steps
  - Type of connections: ONE, ZERO, LIMIT\_ONCE, LIMIT\_ALWAYS
- Adds placeholders to avoid vertex and swap conflicts by miming the wait action.
- Usually  $A^*$  is used



# Increasing Cost Tree Search (ICTS)

- It was proposed in 2013 by G. Sharon, R. Stern, M. Goldenberg and A. Felner and it is optimal [7]
- Similarly to CBS, ICTS is divided in two searches: high-level and low-level
- It uses an Increasing Cost Tree
- For each new level,  $k$  new nodes are created
- The search continues until a feasible solution is found




# ICTS – Low-Level Search


- The low-level search is implemented using Multi-value Decision Diagrams (MDDs) [8]
- An MDD contains all the paths for an agent going from its initial position to their goal with a certain cost



Cost 3





Cost 4




Cost 5

# ICTS – Low-Level Search


- We can merge the MDDs of different agents to check for possible solutions
- The branches that have conflicts are removed



Red agent



Blue agent



Merged MDD

# Constraint Programming (CP) [9]

- MAPF has been proven to be NP-Hard [10] → it can be reduced to SAT and MILP
- Constraint programming is a mathematical modeling paradigm in which some constraints are placed over some variables.
- Some constraints are:
  - Agents must be only on one vertex at each time step;
  - A node can be occupied by at most one agent at a time;
  - Agents start from their initial position and must be on their arrival position at the end;
  - Agents must move along edges.

# Constraint Programming – Implementation

- Started from the work of Bartak et al using Picat
- Moved to IBM's CPLEX for performance
- Decision variables:

$X[n\_steps][n\_nodes][n\_agents]$  movement[n\_agents][n\_steps]  
 $goal\_points[n\_steps][n\_nodes][n\_agents]$  edges[n\_agents][n\_steps]

- An agent  $a$  can be only on one node  $n$  at a time  $s$ :

$$\forall s \in S, \forall a \in A, \sum_{n \in N} X[s][n][a] = 1$$

- Python:

```
for s in steps:
    for a in agents:
        m.add_constraint((m.sum(x[(s, n, a)] for n in nodes) == 1))
```

- C++:

```
FOREACH(s, steps) {
    FOREACH(a, agents) {
        IloExpr expr(env);
        FOREACH(n, nodes) {
            expr += x[s][n][a];
        }
        model.add(x: expr <= 1);
    }
}
```



# Constraint Programming – Constraints

- Examples of constraints are:
  - Agents cannot be on more than one node each time step
  - A node cannot be occupied by more than one agent at time
  - An agent must occupy a neighbor of the node it is on at time  $t + 1$  or stay on the same node
  - Agents start on their initial positions and end on their final positions
  - At a certain time, an edge cannot be used in more than one direction
  - The movement cost of an agent at a given time is the cost of the edge it is traversing
  - The agent must go through all the goals and only once before reaching the final position


# Extended ICTS [11]

- Standard MAPF considers edges with unit costs
- Assumptions on agents' movements:
  - Wait on the center of the node
  - Moves in a straight line
  - Collision is the overlapping in an instant of time
- Two problem:
  - Partial time overlap conflict detection
    - Detect conflicts
  - Partial time overlap successor generation
    - Generate successive states


# Extended ICTS

	1	2	3
A			G
B	S		

(a)




(b)



(c)

# Extended ICTS – Optimal

- In ICTS, the MDD had one root and one sink
- The next child is not obtained with an increment of 1, but we need to set an increment value  $\delta$ 
  - If  $\delta$  is small --> the depth of the ICT may become too big
  - If  $\delta$  is large --> search is reduced, but the solution may not be optimal
- The ICT nodes now contain intervals:
  - Lower bound is the solution minimum
  - Higher bound is the solution maximum
- The low-level is changed from a satisfactory problem to an optimal one: find the solution with the minimum cost in the interval



# Extended ICTS – Heuristics

- $\epsilon$ -ICTS: considers the low-level a satisfactory problem
  - This allows the algorithm to find a solution that is bounded sub-optimal
- $w$ -ICTS: we can bound the sub-optimality for the generation of the next step to values of  $\delta$  by adding a weight value  $w$

# CBS Heuristics


- CBS has been the start of the show with many improvements:
  - Bypassing conflicts [12]
  - Prioritizing conflicts [13]
  - Symmetry reasoning [14]
- And also a number of different heuristics:
  - ECBS [15]
  - EECBS [16]
  - EEEEECBS (Saccon et al., 2025)
  - EEEEEEEEEEEEEEEEECBS (Saccon et al., 2030).

# CBS Heuristics

- CBS has been the start of the show with many improvements:
  - Bypassing conflicts [12]
  - Prioritizing conflicts [13]
  - Symmetry reasoning [14]
- And also a number of different heuristics:
  - ECBS [15]
  - EECBS [16]

# CBS Heuristics

- Bypassing conflicts:
  - Do not split the node every time a conflict is found
  - When analyzing a conflict, modify the agents' path
  - If the cost of the new solution is the same as before and the number of conflicts is reduced --> do not split, but substitute
- Prioritizing conflicts:
  - A conflict is cardinal iff by solving the cost of both child CT nodes increases
  - A conflict is semi-cardinal iff the cost of only one child increases
  - A conflict is non-cardinal iff neither children's cost increased
- Symmetry reasoning





# Enhanced CBS (ECBS)

- Instead of using A\* uses Focal search [17]
  - Bounded suboptimal algorithm
  - Uses OPEN and FOCAL sets
  - FOCAL contains all those nodes that have a weights suboptimal cost
  - The values in FOCAL are sorted using a function to estimate the cost-to-go
- ECBS implements focal search both for the low-level search and the high-level search:
  - The low-level is not sped up --> given an agent and a CT node, it returns
    - the path that minimizes the number of conflicts with other agents
    - the cost of the shortest path
  - The high-level search gets the costs of the shortest paths and can use focal search to speed up the analyses of the tree

# Explicit Estimation Search

- Two main problems with ECBS high-level search:
  - It considers only the cost-to-go --> solution cost may be greater than the sub-optimality bound
  - At each time, the number of CT nodes with a similar cost is large --> FOCAL is rarely emptied
- The world is full of heuristic searches!
- Explicit Estimation Search (EES) [18] is a bounded-suboptimal search algorithm --> uses one more heuristic to overcome said focal behavior

# Explicit Estimation Search (EES)

- EES is a bounded-suboptimal search algorithm --> uses one more heuristic to overcome said focal behavior
- It uses  $\hat{h}$  and  $\hat{d}$  to estimate the cost-to-go and the distance-to-go
- It keeps track of:
  - $best_f$ , the node minimizing  $f(n) = g(n) + h(n)$  from the FOCAL list
  - $best_{\hat{f}}$ , the lowest predicted solution cost
  - $best_{\hat{d}}$ , the node between the  $w$  admissible ones that appears closer to the target
- The node to explore is chosen based on
  1.  $\hat{f}(best_{\hat{d}}) \leq w \cdot f(best_f) \rightarrow best_{\hat{d}}$  --> chose the node nearest to the goal
  2.  $\hat{f}(best_{\hat{f}}) \leq w \cdot f(best_f) \rightarrow best_{\hat{f}}$  --> chose the node with the best path
  3.  $best_f$  --> trust  $A^*$

# Explicit Estimation CBS (EECBS)

- EECBS improves on ECBS by using EES on the *high-level* search
- It maintains 3 lists of CT nodes:
  - CLEANUP: regular list of  $A^*$  sorted by the lower bound
  - OPEN: regular list of  $A^*$  sorted by a *potentially inadmissible* function
  - FOCAL: nodes with cost bounded by  $w$  sorted by the distance-to-go
- The choice of the node to expand is similar to EES:
  - $cost(best_{h_c}) \leq w \cdot lb(best_{lb}) \rightarrow best_{h_c}$
  - $cost(best_{\hat{f}}) \leq w \cdot lb(best_{lb}) \rightarrow best_{\hat{f}}$
  - $best_{lb}$

# Anytime Solvers

- We have seen that:
  - optimal algorithms do not scale well for the problem, but
  - sub-optimal algorithms may return a solution that is too inadequate
- Here come anytime solvers!
- The idea is:
  - return a sub-optimal solution in the shortest amount of time possible;
  - refine said solution in the remaining time available
- Many algorithms focus on intersections --> nodes with two or more neighbors
- We will see:
  - MAPF-LNS [19]
  - $X^*$  [21]

# MAPF-LNS

- The algorithm uses Large-Neighborhood Search (LNS) [20]
  - The idea is to take a solution, remove an area, consider what remains as good and replan only on the sub-area which is a sub-problem of the initial
- It starts by finding an initial sub-optimal solution with either
  - EECBS, PP, or heuristics on PP
- The important aspect is how to extract a neighborhood
  - Agent-based
  - Map-based
  - Random-based

# MAPF-LNS


- Agent-based neighborhood:
  1. Extract those agents that are not following the shortest path they could
  2. For each, compute a shorter random path
  3. Find agents that are colliding
  4. Replan groups of agents
- Map-based neighborhood:
  1. Identify the intersections --> higher probability of collision
  2. Identify agents moving through intersection, or in the area
  3. Change order in which agents pass through intersection
- Random neighborhood:
  1. Randomly choose N agents to replan for
  2. Replan

# X\*


- They introduce a concept called *window*
  - Identify agents and states around a conflict and repair the conflict
  - Each window has a successor, which basically is a larger window --> windows can grow
  - Two windows can also be merged together
- How does it work?
  1. Plan each agent individually
  2. Then starting by time  $t_0$  it looks for conflicts
  3. For each conflict it creates a window and tries to solve it locally to the window
  4. The fixes are optimal within the window
- By enlarging windows and merging them, it can produce an optimal solution
- By preventing changes to following windows, it can produce a sub-optimal solution




# $X^*$ – Example




(a) Individually planned paths for each agent from  $s$  to  $g$  are used to form a global path. An agent-agent collision occurs in the path between  $a$  and  $b$  at  $t = 3$ .



(b) Collision between  $a$  and  $b$  is repaired by jointly planning inside  $w^{ab}$ . The global path is now guaranteed to be valid, but not guaranteed to be optimal.




(c)  $w^{ab}$  is grown and a new repair is generated for  $a$  and  $b$ . The window does not yet encapsulate the search from  $s^{ab}$  and  $g^{ab}$ , so the repaired global path is not yet guaranteed to be optimal.



(d)  $w^{ab}$  is grown and a new repair is generated. The repair search is from  $s^{ab}$  to  $g^{ab}$  and unimpeded by  $w^{ab}$ , thus allowing  $w^{ab}$  to be removed and the global path returned as optimal.

# X\*— Merging




(a) Individually planned paths for each agent from  $s$  to  $g$  are used to form a global path. An agent-agent collision occurs between  $a$  and  $b$  at  $t = 1$ .



(b) Collision between  $a$  and  $b$  is repaired by jointly planning inside  $w^{ab}$ . The repair creates a collision between  $a$  and  $c$  at  $t = 7$ .




(c) Collision between  $a$  and  $c$  and is repaired by jointly planning inside  $w^{ac}$ . No collisions exist, thus producing a valid global path.



(d) All windows are grown in order to improve repair quality.




(e) As they overlap in agent set and states,  $w^{ab}$  and  $w^{ac}$  are merged to form  $w^{abc}$ , and a new repair is generated and inserted into the global path.




(f)  $w^{abc}$  is repeatedly grown and searched until the search of  $w^{abc}$  takes place from  $s^{abc}$  to  $g^{abc}$  unimpeded, thus allowing  $w^{abc}$  to be removed and the global path returned as optimal.

# X\* – Example




$a$  path:  $\leftarrow \leftarrow \leftarrow \leftarrow \leftarrow \uparrow$   
 $b$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \rightarrow$   
 $c$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \uparrow \uparrow \uparrow \uparrow$

(a) Individually planned paths for each agent from  $s$  to  $g$  are used to form a global path. An agent-agent collision occurs in the path between  $a$  and  $b$  between  $t = 3$  and  $t = 4$ . Walls are depicted by thick black lines.




$a$  path:  $\leftarrow \leftarrow \leftarrow \leftarrow \leftarrow \uparrow$   
 $b$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \rightarrow$   
 $c$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \uparrow \uparrow \uparrow \uparrow$

(b) Collision between  $a$  and  $b$  is repaired by jointly planning inside  $w^{ab}$ .  $b$  now side steps into the slot to allow  $a$  to pass, but this repair causes a collision with  $c$  at  $t = 9$ . The region of the paths repaired by  $w^{ab}$  is surrounded by dashed lines.



$a$  path:  $\leftarrow \leftarrow \leftarrow \leftarrow \leftarrow \uparrow$   
 $b$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \rightarrow$   
 $c$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \uparrow \uparrow \uparrow \uparrow$



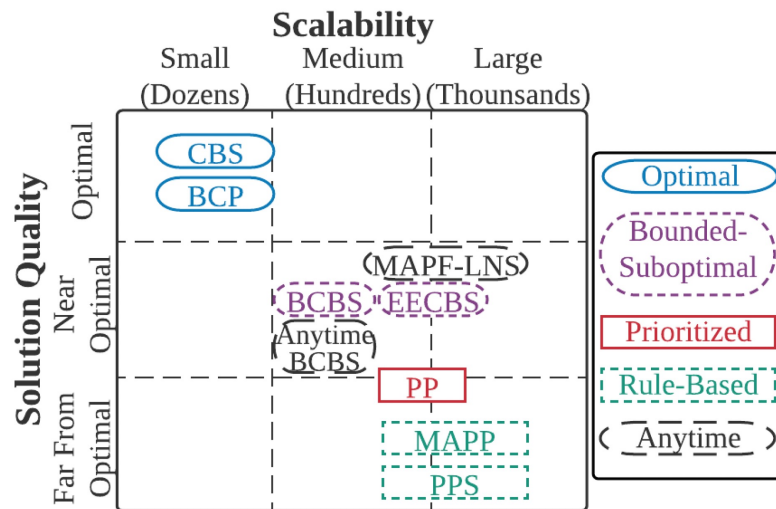
$a$  path:  $\leftarrow \leftarrow \leftarrow \leftarrow \leftarrow \uparrow$   
 $b$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \rightarrow$   
 $c$  path:  $\rightarrow \rightarrow \rightarrow \rightarrow \rightarrow \uparrow \uparrow \uparrow \uparrow$

# MAPF – Variants

- Different types of conflicts
- Different behaviors of agents when reaching goal
- MAPF with agents of different sizes
- Lifelong MAPF
- MAPF with non discrete time
- Multi-Objective MAPF
- MAPF in combination with task planning

# Recap

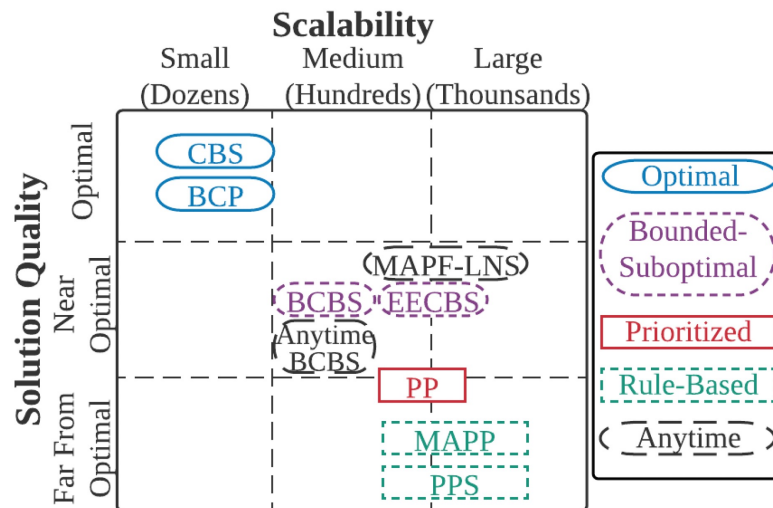
- Optimal:
  - Enhanced A\* [2] (complete with correct heuristic)
- Complete and optimal:
  - CBS [6]
  - ICTS [7]
  - Constraint/logic programming [9]
- Fast:
  - PP [4]
- Suboptimal:
  - $\epsilon$ -ICTS and  $w$ -ICTS
  - CBS with improvements
  - ECBS
  - EECBS
- Anytime solvers:
  - MAPF-LNS
  - $X^*$



# Recap

- Optimal:
  - Enhanced A\* [2] (complete with correct heuristic)
- Complete and optimal:
  - CBS [6]
  - ICTS [7]
  - Constraint/logic programming [9]
- Fast:
  - PP [4]
- Suboptimal:
  - $\epsilon$ -ICTS and  $w$ -ICTS
  - CBS with improvements
  - ECBS
  - EECBS
- Anytime solvers:
  - MAPF-LNS
  - $X^*$

[MAPF Bible](#)



# References

- [1] Stern, Roni, et al. "Multi-agent pathfinding: Definitions, variants, and benchmarks." *Proceedings of the International Symposium on Combinatorial Search*. Vol. 10. No. 1. 2019.
- [2] <https://www.leagueofrobotrunners.org/>
- [3] Standley, T. (2010). Finding Optimal Solutions to Cooperative Pathfinding Problems. *Proceedings of the AAAI Conference on Artificial Intelligence*, 24(1), 173-178. <https://doi.org/10.1609/aaai.v24i1.7564>
- [4] Silver, D. (2021). Cooperative Pathfinding. *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, 1(1), 117-122. <https://doi.org/10.1609/aiide.v1i1.18726>
- [5] Ma, H., Harabor, D., Stuckey, P. J., Li, J., & Koenig, S. (2019). Searching with Consistent Prioritization for Multi-Agent Path Finding. *Proceedings of the AAAI Conference on Artificial Intelligence*, 33(01), 7643-7650. <https://doi.org/10.1609/aaai.v33i01.33017643>
- [6] Sharon, Guni, et al. "Conflict-based search for optimal multi-agent pathfinding." *Artificial Intelligence* 219 (2015): 40-66.
- [7] Sharon, Guni, et al. "The increasing cost tree search for optimal multi-agent pathfinding." *Artificial intelligence* 195 (2013): 470-495.
- [8] A. Srinivasan, T. Ham, S. Malik and R. K. Brayton, "Algorithms for discrete function manipulation," *1990 IEEE International Conference on Computer-Aided Design. Digest of Technical Papers*, Santa Clara, CA, USA, 1990, pp. 92-95, doi: 10.1109/ICCAD.1990.129849.
- [9] R. Barták, N. -F. Zhou, R. Stern, E. Boyarski and P. Surynek, "Modeling and Solving the Multi-agent Pathfinding Problem in Picat," *2017 IEEE 29th International Conference on Tools with Artificial Intelligence (ICTAI)*, Boston, MA, USA, 2017, pp. 959-966, doi: 10.1109/ICTAI.2017.00147.
- [10] Yu, J., & LaValle, S. (2013). Structure and Intractability of Optimal Multi-Robot Path Planning on Graphs. *Proceedings of the AAAI Conference on Artificial Intelligence*, 27(1), 1443-1449. <https://doi.org/10.1609/aaai.v27i1.8541>
- [11] Walker, Thayne T., Nathan R. Sturtevant, and Ariel Felner. "Extended Increasing Cost Tree Search for Non-Unit Cost Domains." *IJCAI*. 2018.

# References

- [12] Boyrasky, Eli, et al. "Don't split, try to work it out: Bypassing conflicts in multi-agent pathfinding." *Proceedings of the International Conference on Automated Planning and Scheduling*. Vol. 25. 2015.
- [13] Boyarski, Eli, et al. "Icbs: The improved conflict-based search algorithm for multi-agent pathfinding." *Proceedings of the International Symposium on Combinatorial Search*. Vol. 6. No. 1. 2015.
- [14] Li, Jiaoyang, et al. "New techniques for pairwise symmetry breaking in multi-agent path finding." *Proceedings of the International Conference on Automated Planning and Scheduling*. Vol. 30. 2020.
- [15] Barer, Max, et al. "Suboptimal variants of the conflict-based search algorithm for the multi-agent pathfinding problem." *Proceedings of the International Symposium on Combinatorial Search*. Vol. 5. No. 1. 2014.
- [16] Li, J., Ruml, W., & Koenig, S. (2021). EECBS: A Bounded-Suboptimal Search for Multi-Agent Path Finding. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(14), 12353-12362. <https://doi.org/10.1609/aaai.v35i14.17466>
- [17] Pearl, Judea, and Jin H. Kim. "Studies in semi-admissible heuristics." *IEEE transactions on pattern analysis and machine intelligence* 4 (1982): 392-399.
- [18] Thayer, Jordan Tyler, and Wheeler Ruml. "Bounded suboptimal search: A direct approach using inadmissible estimates." *IJCAI*. Vol. 2011. 2011.
- [19] Li, Jiaoyang, et al. "Anytime multi-agent path finding via large neighborhood search." *International Joint Conference on Artificial Intelligence 2021*. Association for the Advancement of Artificial Intelligence (AAAI), 2021.
- [20] Shaw, Paul. "Using constraint programming and local search methods to solve vehicle routing problems." *International conference on principles and practice of constraint programming*. Berlin, Heidelberg: Springer Berlin Heidelberg, 1998
- [21] Vedder, Kyle, and Joydeep Biswas. "X\*: Anytime multi-agent path finding for sparse domains using window-based iterative repairs." *Artificial Intelligence* 291 (2021): 103417.



